

PyTorch Tutorial

CSE455

Kiana Ehsani





What is PyTorch?

PYTORCH



Caffe



PyTorch takes care of everything



Andrej Karpathy ✓

@karpathy



I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

11:56 AM - May 26, 2017



1,512



418 people are talking about this



Installation

<https://pytorch.org/>


https://github.com/ehsanik/pytorch_installation

An example

MNIST classifier

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9 9


Data & Labels




0
1
2
3
4
5
6
7
8
9

Data normalization


original data



zero-centered data



normalized data



Modules, Layers and more :)

All networks need to be a child class of `nn.Module`

1. `Conv1d/2d`
2. `MaxPool1d/2d`
3. `MaxUnpool1d/2d`
 - a. `MaxUnpool1d` takes in as input the output of `MaxPool1d` including the indices of the maximal values and computes a partial inverse in which all non-maximal values are set to zero.
4. `Linear`
5. `Upsample`
6. And tons of others


Initialization

1. All Zero -> **constant_**
 - a. PLEASE DON'T!
2. Identity -> **eye_**
3. Random
 - a. Normal distribution -> **normal_**
 - b. Uniform distribution -> **uniform_**
 - c. Xavier
 - d. Kaiming

ALL YOU NEED IS A GOOD INIT





Dmytro Mishkin, Jiri Matas

Center for Machine Perception
Czech Technical University in Prague
Czech Republic {mishkdmy, matas}@cmp.felk.cvut.cz





Non-linear activations

1. What does it mean?
 - a. Neuron is firing!
 - b. Binary! Either 0 or 1
2. What do we have?
 - a. ReLU
 - b. LeakyReLU
 - c. Sigmoid
 - d. Tanh
 - e. Softmax
 - f. And tons of others



Normalizations

1. More stable
2. Reduces the amount by which the hidden unit values shift around




Loss functions

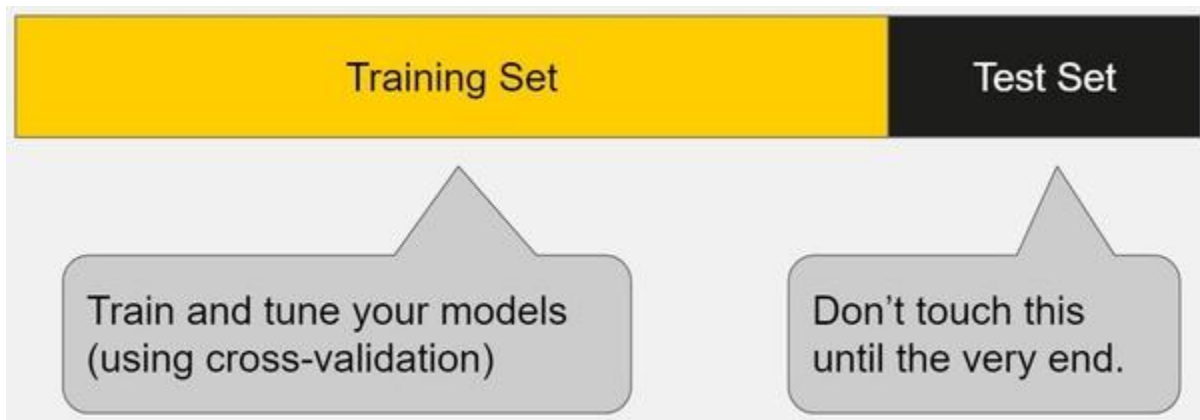
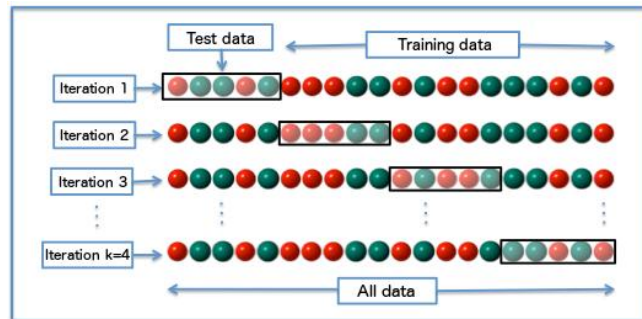
1. L1Loss
2. MSELoss
3. CrossEntropyLoss
4. BCELoss
5. KL Divergence Loss
 - a. Useful when direct regression over a distribution
6. And tons of others
7. Implement your own loss function

Optimizers

1. SGD
2. Adagrad
3. Adam
4. Adadelta
5. And tons of others





Train, Test, Val (Dev)



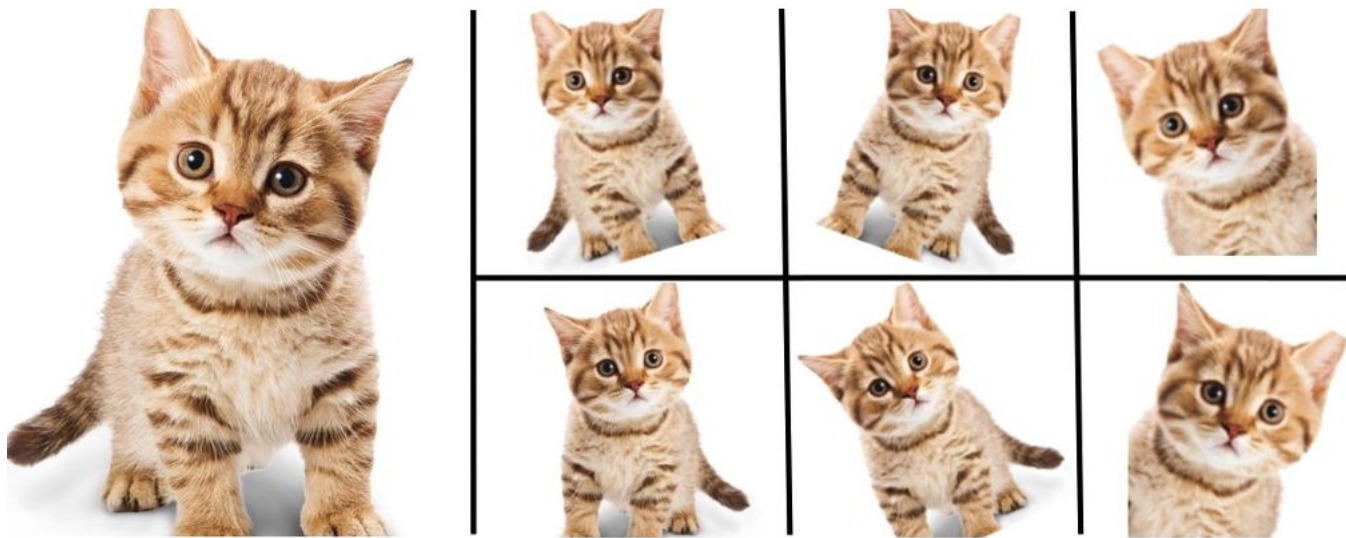
Issues you might face

1. Overfitting
2. Getting stuck in local minima
3. Diverging

Overfitting



Data transformations




Enlarge your Dataset


Data transformations

1. Grayscale
2. ColorJitter
3. LinearTransformation
4. Pad
5. RandomAffine
6. RandomCrop
7. RandomHorizontalFlip
8. RandomRotation


Be careful! It depends on the dataset.



Regularization: Dropout



(a) Standard Neural Net




(b) After applying dropout.

Dropout vs Batchnorm




Batch normalization



Dropout



Diverging and local minima




Project Ideas

Datasets:

<https://docs.google.com/spreadsheets/d/1pXCrtVdPYOJthsBaMyRLQz6jLCor55vbcoKgetSvexU/edit#gid=0>

Ideas? Brain Storming!

Backstage



A graph is created on the fly

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))
```



$$b = w_1 * a$$

$$c = w_2 * a$$

$$d = (w_3 * b) + (w_4 * c)$$

$$L = f(d)$$

